

Broom: Safe Region-based Memory Management for BigData Processing Systems

with G Ramalingam, Kapil Vaswani, Dimitrios Vytiniotis,
Jana Gicheva, Michael Isard, Steve Hand & Ionel Gog



Microsoft

In this talk ...

can be framed and solved

How **Memory Safety** as Type
Safety in Broom environment



Memory Safety as Type Safety in Broom

Broom

Broom

Optimized software stack for big data pipelines

DryadLinq Naiad CScale

data processing pipelines for big data

compute count on
(select .. from ...)

- large
- semi-structured
- distributed


Broom

```
compute count on  
(select ... from ... )
```

Broom

compute count on
(select ... from ...)

computations can be
staged




Broom

compute mean on
(select ... from ...)

computations can be
staged

... where each stage
is a data parallel
computation



Most big data computations can be represented as
dataflow DAGs

Broom


Optimized software stack for big data pipelines

DryadLinq Naiad CScale

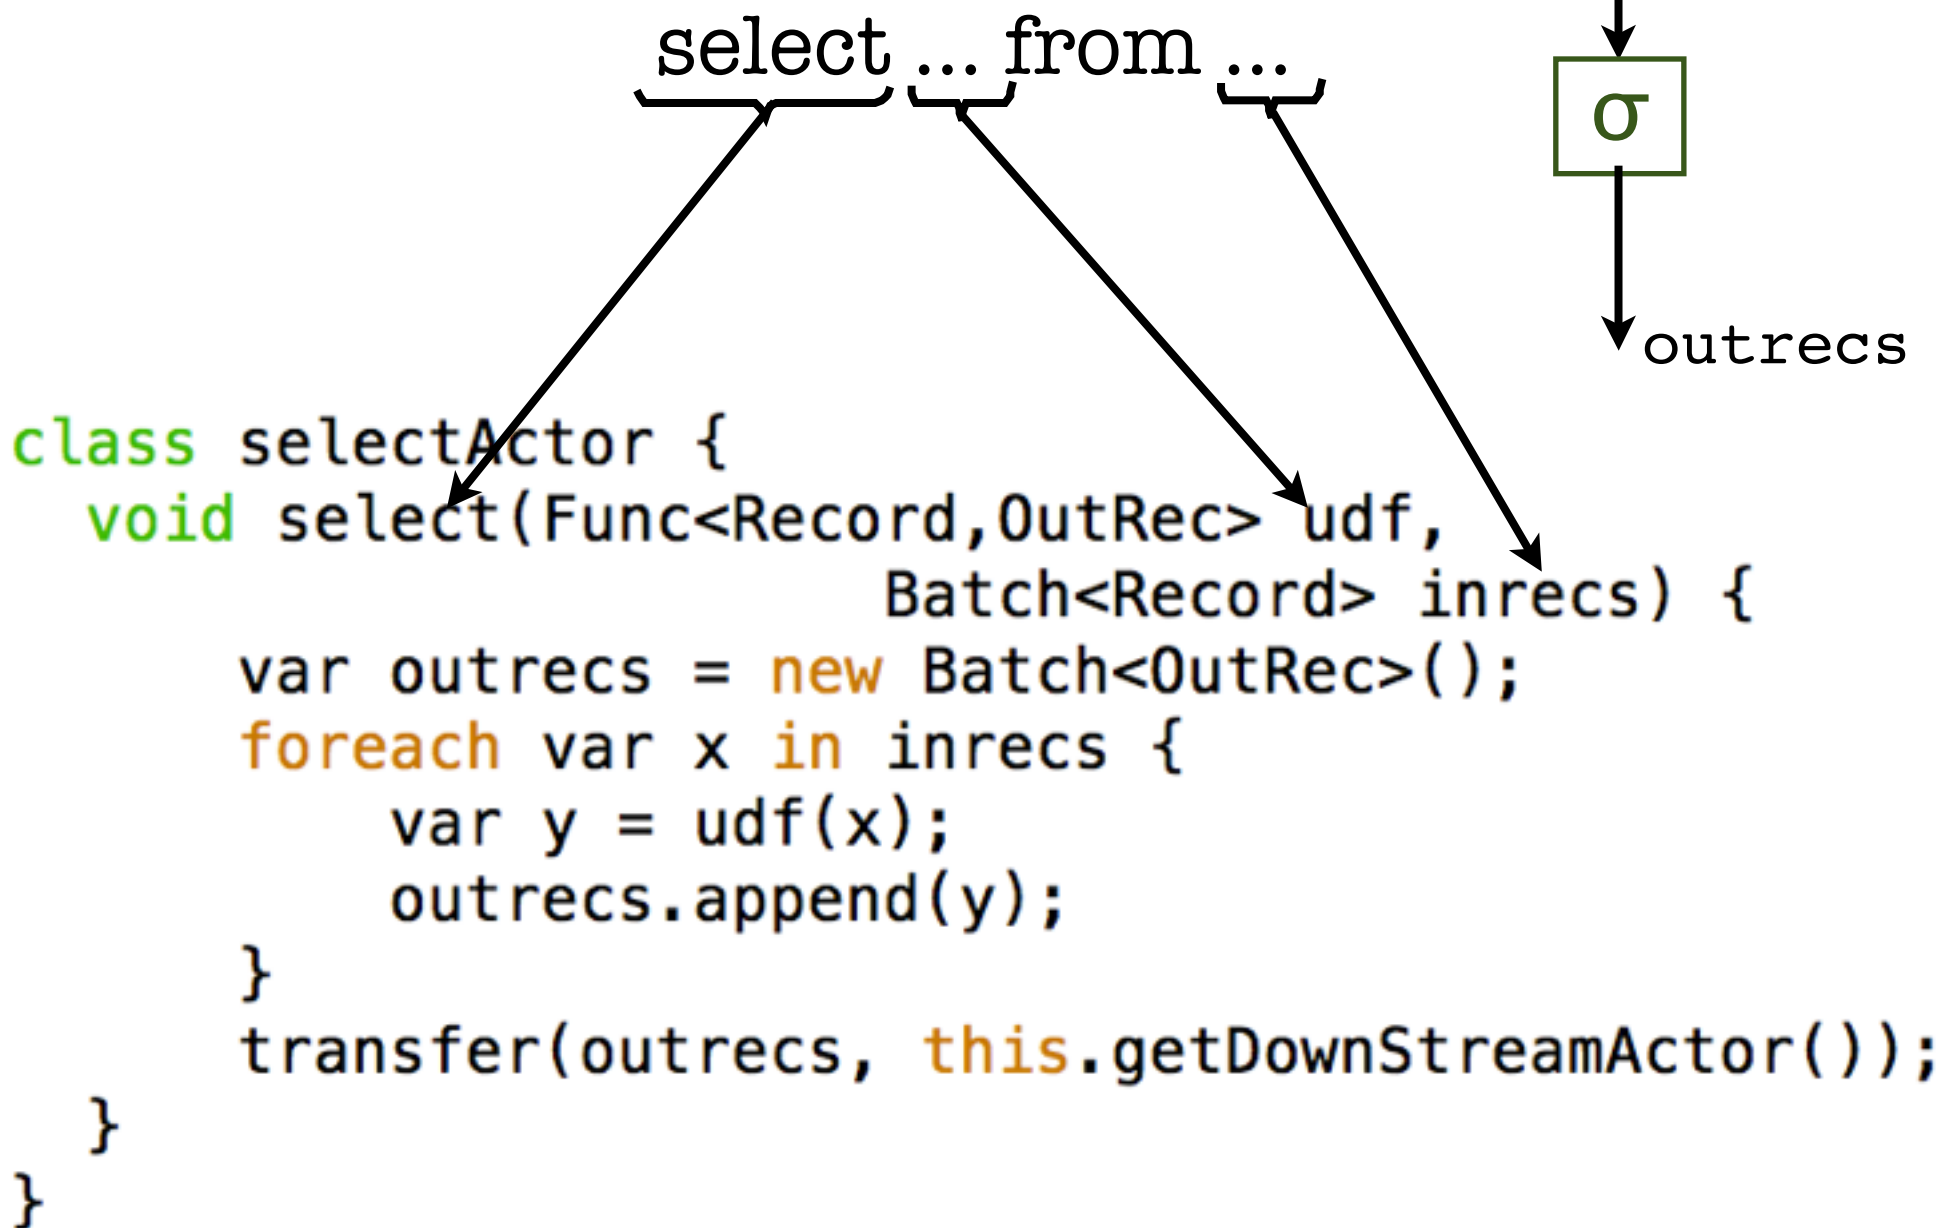
Dryad is a general-purpose distributed execution engine for coarse-grain data-parallel applications. A Dryad application combines computational “vertices” with communication “channels” to form a dataflow graph. Dryad runs the application by executing the vertices of this graph on a set of available computers, communicating as appropriate through files, TCP pipes, and shared-memory FIFOs.

computation vertices communicating
via message-passing channels

Actors



An Example Select Actor σ



Broom

Current software stack for big data pipelines



General purpose:

- ★ Garbage collection
- ★ Scheduling
- ★ I/O management

Not necessarily optimized for
distributed dataflow pipelines

Broom

Optimized software stack for big data pipelines

More performant run-time by exploiting specific characteristics of dataflow programs

Broom

for this talk ...

Memory Optimized software stack for big data pipelines

More **memory** performant run-time by exploiting
specific characteristics of dataflow programs

The aim is to bring down the GC overhead

Broom Memory Management

“The basic concept of regions is very old, first appearing as early as 1967 in Douglas T. Ross's AED Free Storage Package” - wikipedia

Broom adopts a well-known memory management scheme called region-based memory management

“A region is a logical container for objects that obey some memory management discipline”

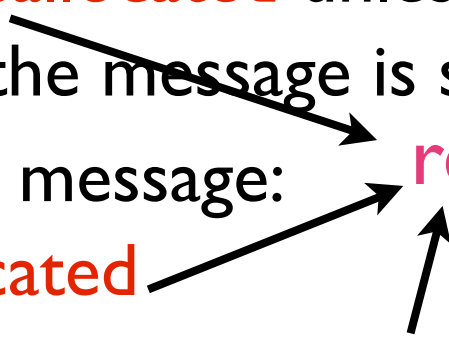
- Hicks, Morrisett et al., Cyclone, ISMM'04

For eg: The heap region, A stack frame

Regions

We can create more specific regions, if we are aware of specific memory behaviour of programs

In Broom:

- ★ Data in memory that needs to be sent as a message to a downstream actor **shouldn't be deallocated** unless the downstream actor is remote and the message is sent.
 - ★ Once an actor finishes *acting* on a message:
 - ★ Incoming **message can be deallocated**
 - ★ Any objects created while in *action* **can be reclaimed**.
- 
- regions!

Broom Regions

After observing the memory behaviour of dataflow programs, we introduced different kinds of regions in Broom:

★ Transferable dynamic regions

can be transferred by an actor to another actor

★ Non-transferable static regions:

★ Actor-scoped

★ Action-scoped

deallocated as soon as action finishes execution

★ Block-scoped / Scratchpad

Static Broom Regions

- ★ Lifetime determined by lexical scope
- ★ Usually hold temporary objects

```
R0 {  
    newAlloc {  
        var x = new LinkedListNode<int>(2);  
        RI {  
            newAlloc {  
                var xs = new LinkedList<int>();  
                xs.addAtHead(x);  
            } xs gets deallocated here  
        }  
    } x gets deallocated here  
}
```

Transferable Broom Regions

★ Dynamic Lifetime

★ First class citizens - can be passed to, and returned from fns

```
var r = new Region<Batch<Record, OutRec>>;  
openAlloc r {  
    var msg = new Batch<Record, OutRec>( )  
    r.set(msg);  
}  
...  
r.transfer(this.getDownStreamActor());
```

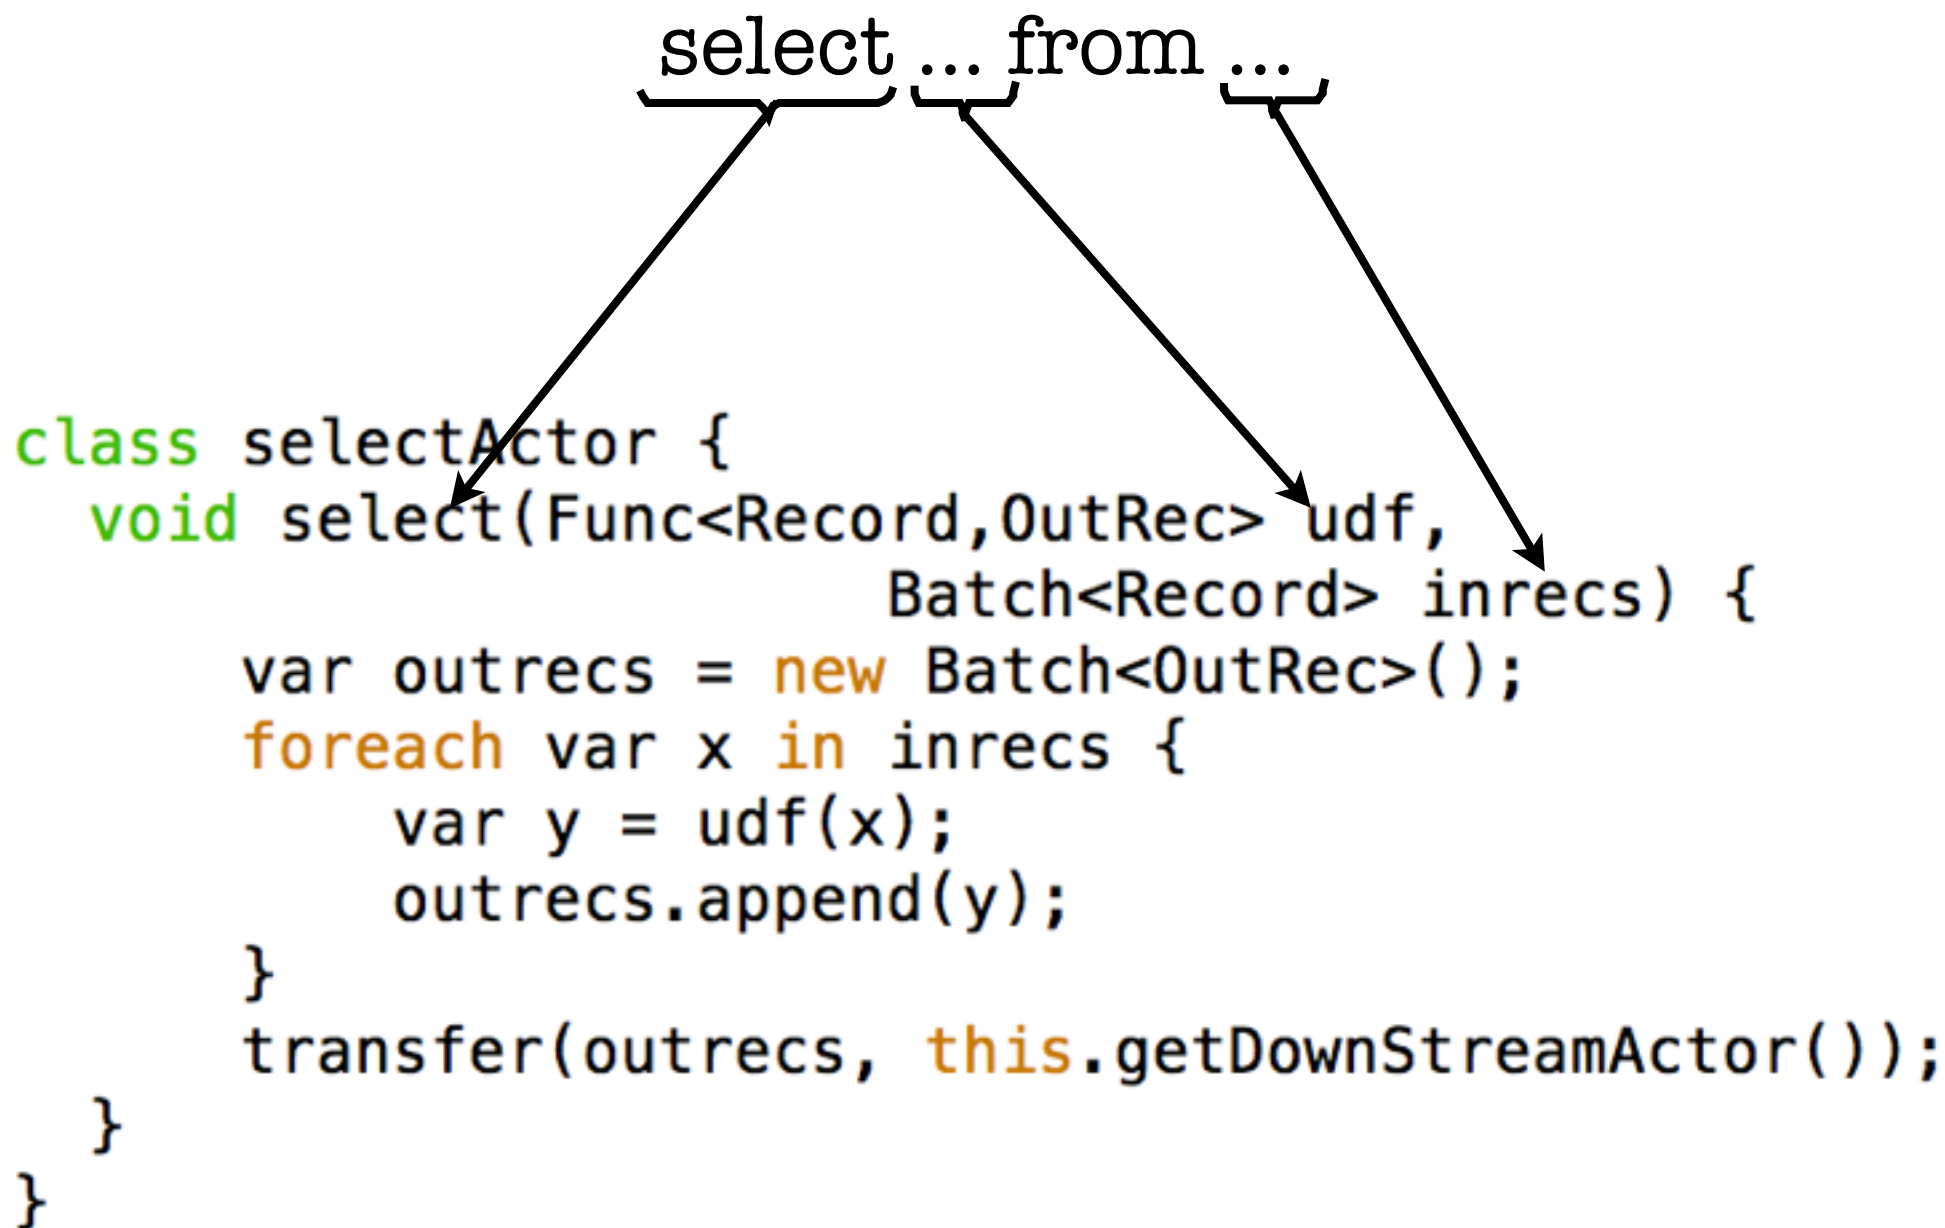
upstream actor

```
void onRecv(Region<Batch<Record, OutRec>> inr) {  
    open inr {  
        ...  
    }  
    inr.giveUp();  
}
```

downstream actor

r

An Example Select Actor σ



Select using Broom Regions

```
void select(Region<Batch<Record>> rin,
            Func<Record, OutRec> udf) {
    → var rtrans = new Region<Batch<OutRecord>>();
    open (rin) {
        var batch = rin.GetData();
        openAlloc (rtrans) {
            rtrans.SetData(new Batch<OutRecord>());
        }
    → newAlloc {
        foreach var x in batch {
            var y = udf(x);
            openAlloc (rtrans) {
                var z = rtrans.SuckObject(y);
                rtrans.GetData().append(z);
            }
        }
    → }
    rin.GiveUp();
    → rtrans.Transfer(this.getDownStreamActor());
}
```


Memory regions have independent life times, and cross-region pointers can become dangling

$R1 \left\{ \begin{array}{l} \text{newAlloc } \{ \\ \dots \\ \text{newAlloc } \{ \\ \dots \\ \} \end{array} \right.$

$R2 \left\{ \begin{array}{l} \text{newAlloc } \{ \\ \dots \\ \} \end{array} \right.$

$R1 \text{ "outlives" } R2: R1 \geq R2$

Therefore, any references from $R1$ to $R2$ are unsafe



It is very easy to shoot oneself in foot.

```

newAlloc {
  var x = new A();
  newAlloc {
    var y = new B();
    x.f = y;
  }
  print (x.f);
}

```

```

var g;
newAlloc {
  var x = new A();
  g = new Func(y => x.m());
}
g();

```

It is easy to create dangling references

... especially with
higher-order functions

```

void foo() {
  var rtrans = new Region<Batch<OutRecord>>();
  openAlloc rtrans {
    ...
  }
  return;
}

```

Space leaks are also possible.


```
newAlloc {  
  var x = new A();  
  newAlloc {  
    var y = new B();  
    x.f = y;  
  }  
  print (x.f);  
}
```

References from live regions to
dead regions

It is easy to create **dangling references**

most of the probable

Can we catch dangling references by analyzing the
code at compile time?

Yes, through a region type system

Tofte, Talpin, Implementing a call-by-value λ -calculus using a stack of
regions, POPL'94

The Region Type System - Intuitions

```
newAlloc {  
  var x = new A();    /* x : A<R0> */  
  newAlloc {  
    var y = new B(); /* y : B<R1> */  
    x.f = y;         /* ∃(R ≥ R0). x.f : B<R> */  
  }  
  print (x.f);  
}
```

Diagram illustrating nested regions: $R0$ contains $R1$, which contains the inner `newAlloc` block.

$x.f = y$ OK?

No!

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B<ρ0> { }`

```
newAlloc<R0> {  
    var x = new B<R0>( );    /* x : B<R0> */  
}
```

```
newAlloc {  
    var x = new A();  
    newAlloc {  
        var y = new B();  
        x.f = y;  
    }  
    print (x.f);  
}
```

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B<ρ0> { }`

`class A {
 B f;
}` \longrightarrow `class A<ρ0> {
 B f;
}`

```
newAlloc {  
    var x = new A();  
    newAlloc {  
        var y = new B();  
        x.f = y;  
    }  
    print (x.f);  
}
```

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B< ρ_0 > { }`

`class A {
 B f;
}` \longrightarrow `class A< ρ_0, ρ_1 > {
 B< ρ_1 > f;
}`

```
newAlloc {  
  var x = new A();  
  newAlloc {  
    var y = new B();  
    x.f = y;  
  }  
  print (x.f);  
}
```

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B< ρ_0 > { }`

`class A {
 B f;
}` \longrightarrow `class A< $\rho_0, \rho_1 \mid \rho_1 \geq \rho_0$ > {
 B< ρ_1 > f;
}`

```
newAlloc {  
  var x = new A();  
  newAlloc {  
    var y = new B();  
    x.f = y;  
  }  
  print (x.f);  
}
```

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B< ρ_0 > { }`

`class A {
 B f;
}` \longrightarrow `class A< $\rho_0, \rho_1 \mid \rho_1 \geq \rho_0$ > {
 B< ρ_1 > f;
}`

```
newAlloc {  
  var x = new A();  
  newAlloc {  
    var y = new B();  
    x.f = y;  
  }  
  print (x.f);  
}
```

$\xrightarrow{\text{newAlloc}\langle R_0 \rangle}$ `new A< $R_0, R?$ >()` but R_1 doesn't exist here!

$\xrightarrow{\text{newAlloc}\langle R_1 \rangle}$ `new B< R_1 >()`

$\underbrace{x.f : B\langle R? \rangle \quad y : B\langle R_1 \rangle}$

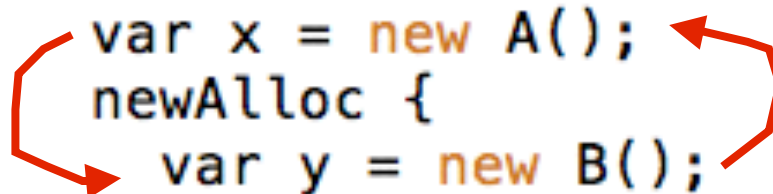
for `x.f=y` to be type-safe, we require $R? = R_1$

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B $\langle\rho_0\rangle$ { }`

`class A {
 B f;
}` \longrightarrow `class A $\langle\rho_0, \rho_1 \mid \rho_1 \geq \rho_0\rangle$ {
 B $\langle\rho_1\rangle$ f;
}`

```
newAlloc {  
  var x = new A();  
  newAlloc {  
    var y = new B();  
    x.f = y;  
  }  
  print (x.f);  
}
```

A diagram illustrating nested memory allocation blocks. Two red arrows originate from the closing curly brace of the inner 'newAlloc' block and point back to the opening curly brace of the outer 'newAlloc' block, indicating that the inner block returns control to the outer one upon completion.

Region Type Checking - Modus Operandi

`class B {}` \longrightarrow `class B< ρ_0 > { }`

`class A {
 B f;
}` \longrightarrow `class A< $\rho_0, \rho_1 \mid \rho_1 \geq \rho_0$ > {
 B< ρ_1 > f;
}`

`newAlloc {
 var y = new B();
 newAlloc {
 var x = new A();
 x.f = y;
 }
 print (x.f);
}`

$\xrightarrow{\text{newAlloc}} \text{newAlloc}<R0>$

$\xrightarrow{\text{new}} \text{new B}<R0>()$

$\xrightarrow{\text{newAlloc}} \text{newAlloc}<R1>$

$\xrightarrow{\text{new}} \text{new A}<R1, R0>()$

$\underbrace{x.f : B<R0> \quad y : B<R0>}$

for `x.f=y` is now perfectly safe

We have looked at examples with non-transferable static regions.

The type system prevented dangling references by only allowing references to longer living regions.

But, does this approach work for transferable regions?

Broom Regions - Example

```
void select(Region<Batch<Record>> rin,
            Func<Record, OutRec> udf) {
    var rtrans = new Region<Batch<OutRecord>>();
    open (rin) {
        var batch = rin.GetData();
        openAlloc (rtrans) {
            rtrans.SetData(new Batch<OutRecord>());
        }
        newAlloc {
            foreach var x in batch {
                var y = udf(x);
                openAlloc (rtrans) {
                    var z = rtrans.SuckObject(y);
                    rtrans.GetData().append(y);
                }
            }
        }
    }
    rin.GiveUp();
    rtrans.Transfer(this.getDownStreamActor());
}
```



y is alloc'ed in static region R

y copied to rtrans

Its lifetime ends
here


possible dangling ptr dereference in
downstream actor

Broom Regions - Example

```
void select(Region<Batch<Record>> rin,
            Func<Record, OutRec> udf) {
    var rtrans = new Region<Batch<OutRecord>>();
    open (rin) {
        var batch = rin.GetData();
        openAlloc (rtrans) {
            rtrans.SetData(new Batch<OutRecord>());
        }
        newAlloc {  newAlloc<R0>
            foreach var x in batch {
                var y = udf(x);  y:OutRec<R0>
                openAlloc (rtrans) {
                    var z = rtrans.SuckObject(y);
                    rtrans.GetData().append(y);
                }
            }
        }
    }
    rin.GiveUp();
    rtrans.Transfer(this.getDownStreamActor());
}
```

If we assume
 $R0 \geq rtrans$
then typecheck
passes!

possible dangling ptr dereference in
downstream actor



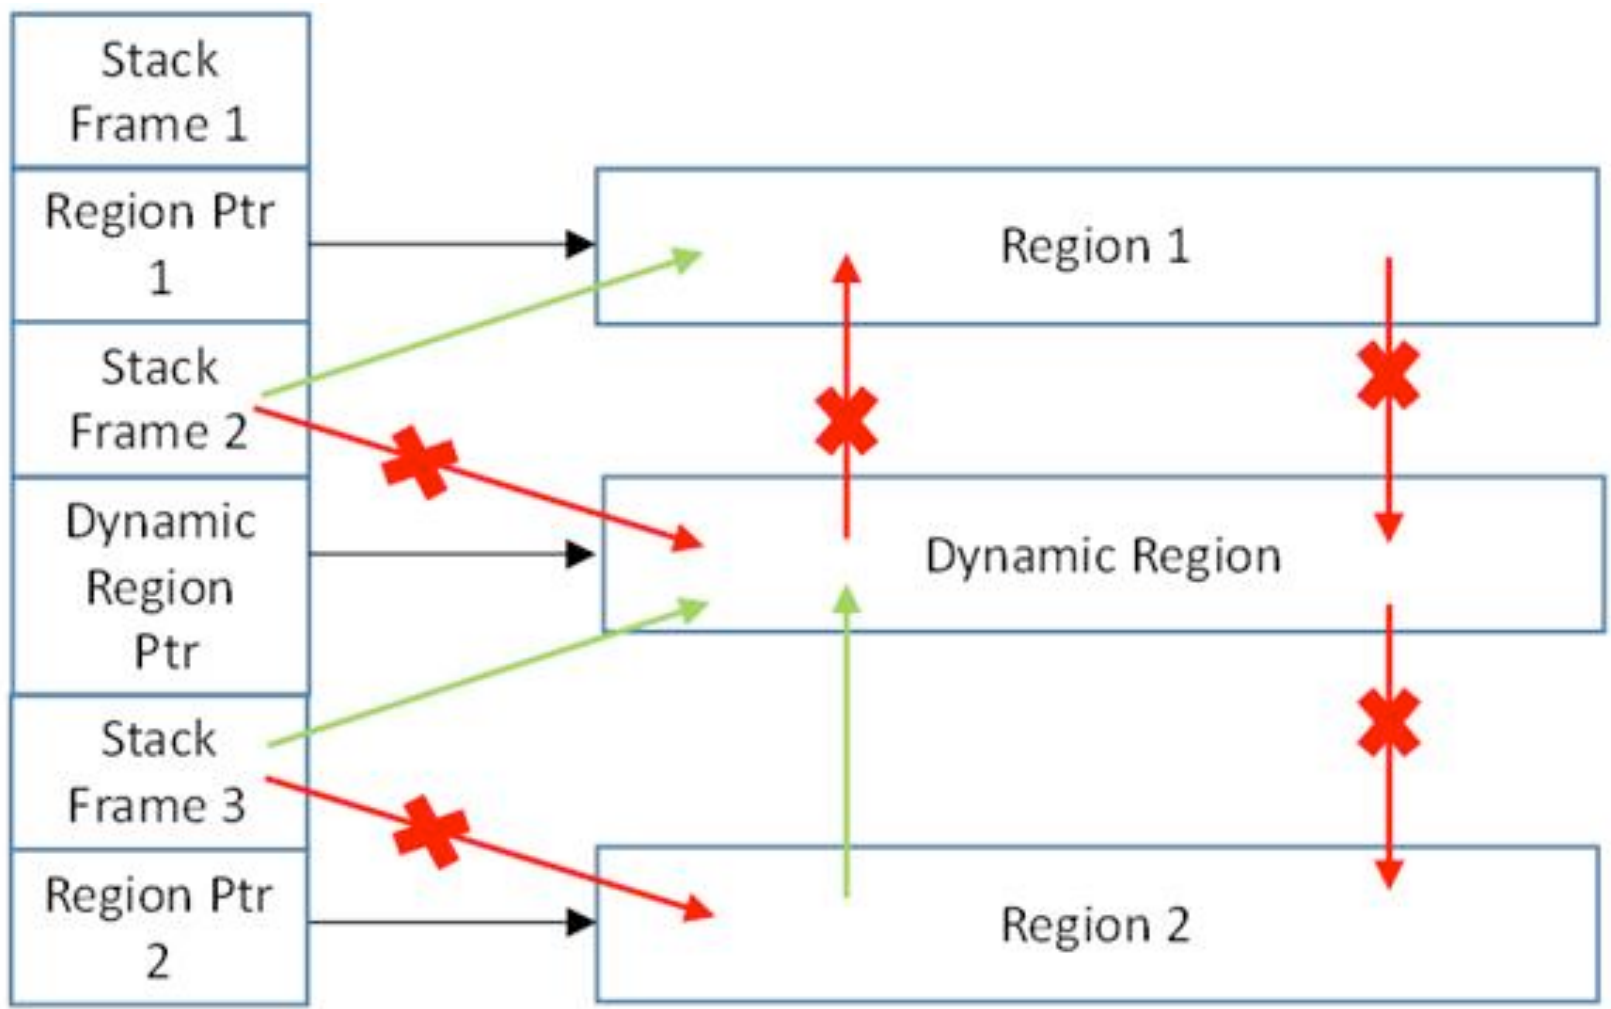

Observation

A sufficient condition to ensure the absence of dangling references from a transferable region is to ensure the absence of *any* references from the transferable region to other regions.





No region outlives a transferable region

stack grows downward



Broom Regions - Example

```
void select(Region<Batch<Record>> rin,
            Func<Record, OutRec> udf) {
    var rtrans = new Region<Batch<OutRecord>>();
    open (rin) {
        var batch = rin.GetData();
        openAlloc (rtrans) {
            rtrans.SetData(new Batch<OutRecord>());
        }
        newAlloc {  newAlloc<R0>
            foreach var x in batch {
                var y = udf(x);  y: OutRec<R0>
                openAlloc (rtrans) {
                    var z = rtrans.SuckObject(y);
                    rtrans.GetData().append(y);
                }
            }
        }
    }
    rin.GiveUp();
    rtrans.Transfer(this.getDownStreamActor());
}
```

Now fails typecheck

possible dangling ptr dereference in
downstream actor

Region Type Inference

Region type checking requires annotating `newAlloc`'s and class names with region names

(eg: `newAlloc<R0>`, `A<R1,R0>`, `B<R0>` etc)

Can these annotations be inferred from plain C# types?

Region Type Inference

Region Type Inference - Intuitions

```
class LinkedListNode<T extends Object> extends
Object {
    T val;
    LinkedListNode<T> prev;
    LinkedListNode<T> next;
    LinkedListNode(T val) {
        super();
        this.val = val;
        this.prev = Null;
        this.head = Null;
    }
}
```

Region Type Inference - Intuitions

```
class LinkedListNode<T> extends Object<ρ1>> extends  
Object<ρ0> {  
    T val;  
    LinkedListNode<T> prev;  
    LinkedListNode<T> next;  
    LinkedListNode(T val) {  
        super();  
        this.val = val;  
        this.prev = Null;  
        this.head = Null;  
    }  
}
```

Region Type Inference - Intuitions

```
class LinkedListNode< $\rho_0, \rho_1$ ><T extends Object< $\rho_1$ >>
extends Object< $\rho_0$ > {
    T val;
    LinkedListNode<T> prev;
    LinkedListNode<T> next;
    LinkedListNode(T val) {
        super();
        this.val = val;
        this.prev = Null;
        this.head = Null;
    }
}
```

Region Type Inference - Intuitions

```
class LinkedListNode< $\rho_0, \rho_1$ ><T> extends Object< $\rho_1$ >>
extends Object< $\rho_0$ > {
    T val;
    LinkedListNode< $\rho_0, \rho_1$ ><T> prev;
    LinkedListNode< $\rho_0, \rho_1$ ><T> next;
    LinkedListNode(T val) {
        super();
        this.val = val;
        this.prev = Null;
        this.head = Null;
    }
}
```

Region Type System

Not covered in today's talk!

- Region type inference -
 - Constraint generation
 - Constraint solving
 - Residual constraints
- Existential typing to assign type to a linked list of transferable regions.
- Why current type system still can't ensure safety with transferable regions.
 - Dynamic solution: throwing exceptions
- InAllocationContext and region-polymorphic types for functions.
- Higher-Order functions.

Region Type System

Open problems

- Region type inference in presence of higher-order functions
- Tracing the failure of constraint satisfaction during type inference back to the erroneous instruction in source program.
- Full region inference: Is it possible to take the code that is written assuming GC, and automatically transform it to a better performing code that uses Broom regions?

Region Type System

Theorems yet to be proved

- **Type safety** - If a broom program has a type under the region type system, then the program does not dereference a pointer into a dead region*
- **Soundness of type inference** - If type inference elaborates (transforms) a program p to program p' , then program p' is well-typed under region type system.

*Conditions Apply

Thank you!