

Mergeable Types

Gowtham Kaki, KC Sivaramakrishnan, Samodya
Abey Siriwardane & Suresh Jagannathan



<https://purduepl.github.io>



<http://ocamlabs.io>



≠



- Replication
- Eventual Consistency

- Serializability
- Linearizability

When **system-level concerns** like replication, consistency etc., affect **application-level design** decisions, programming becomes complicated.

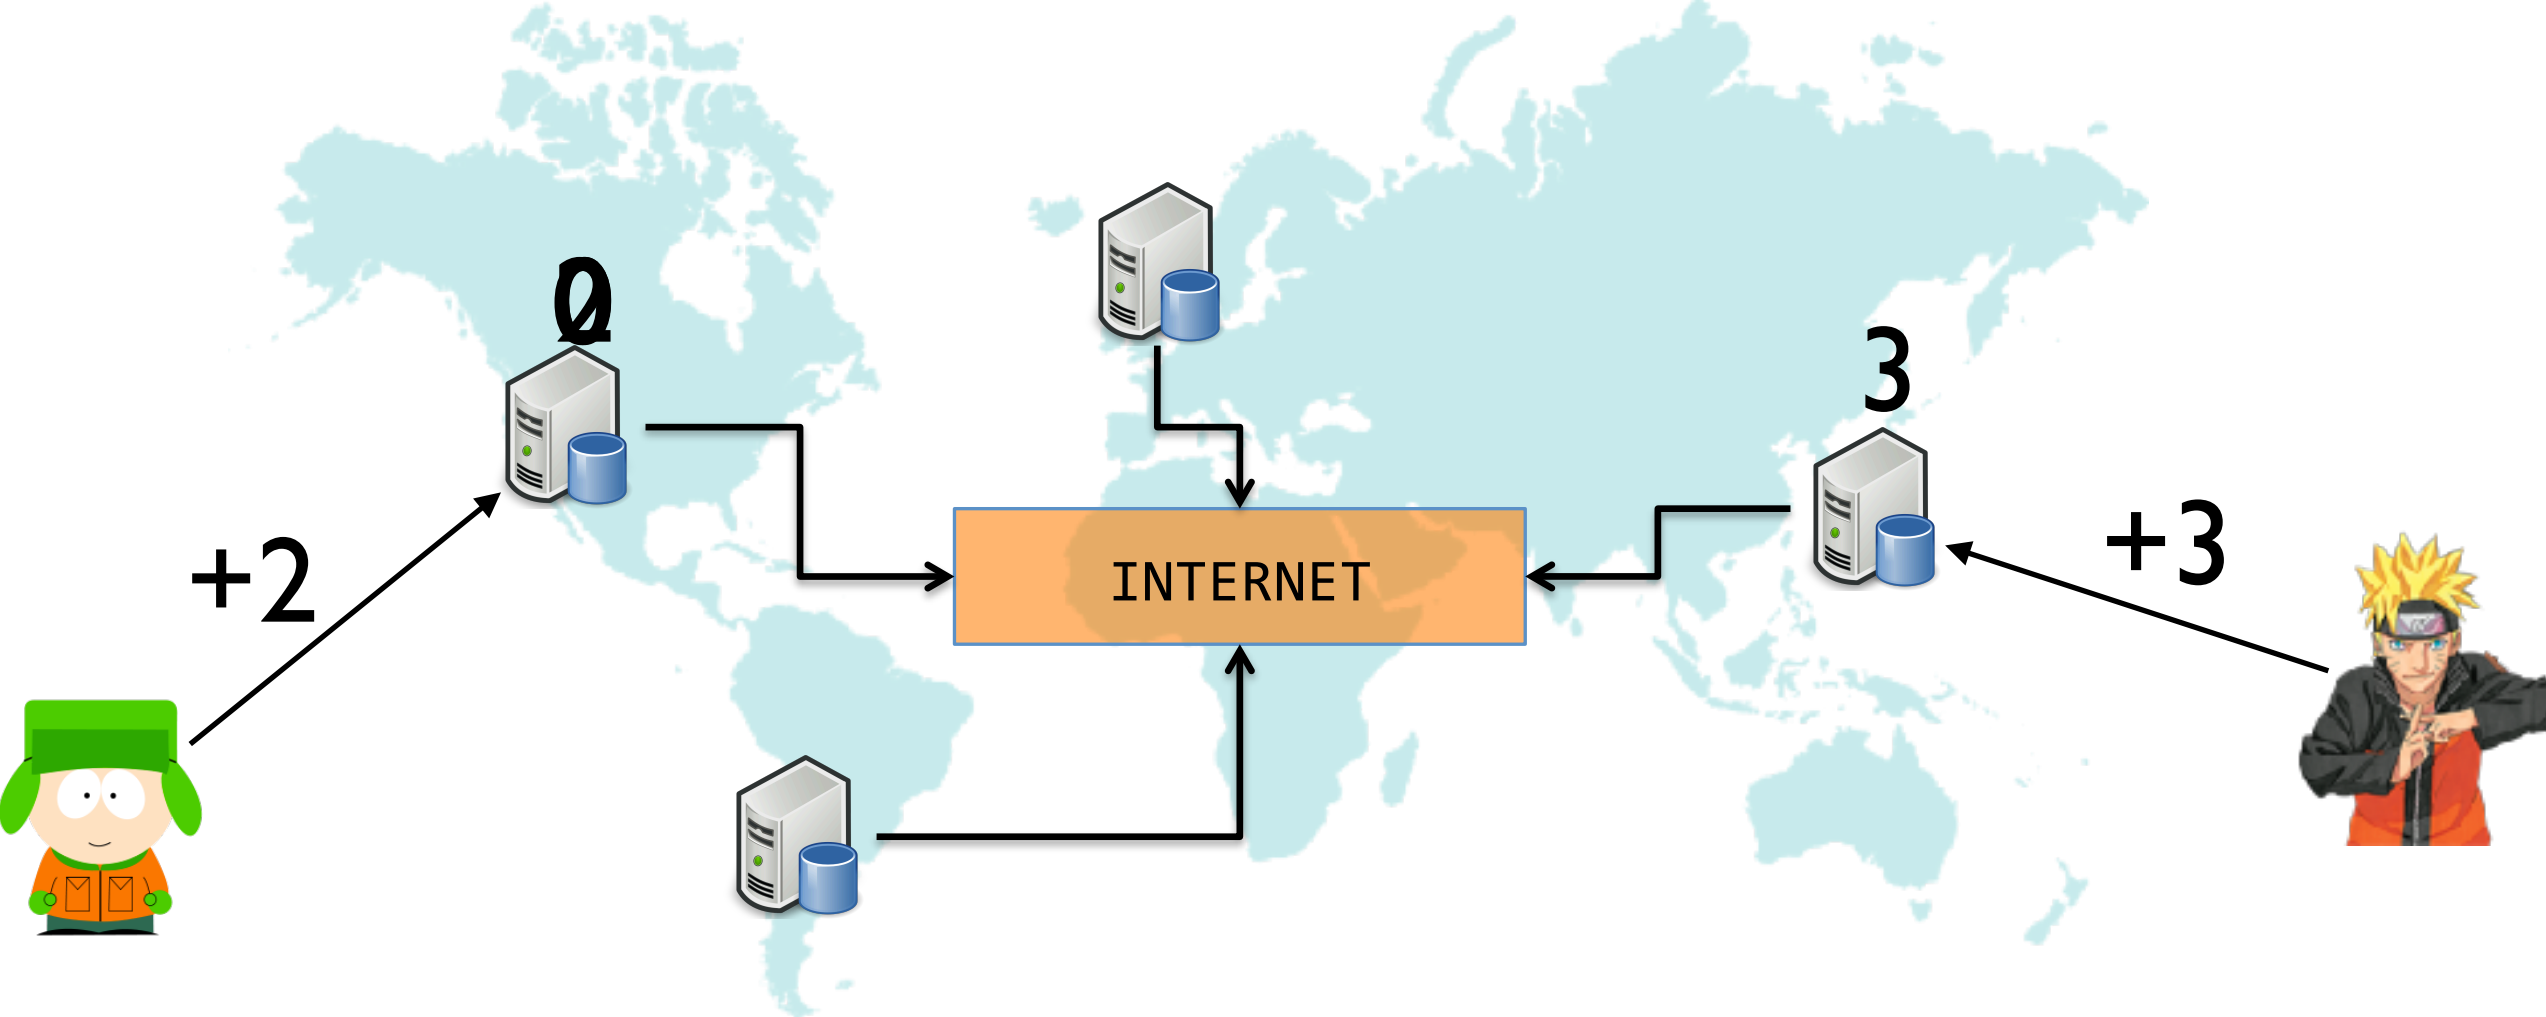
An example...

Monotonic Counter Datatype in OCaml

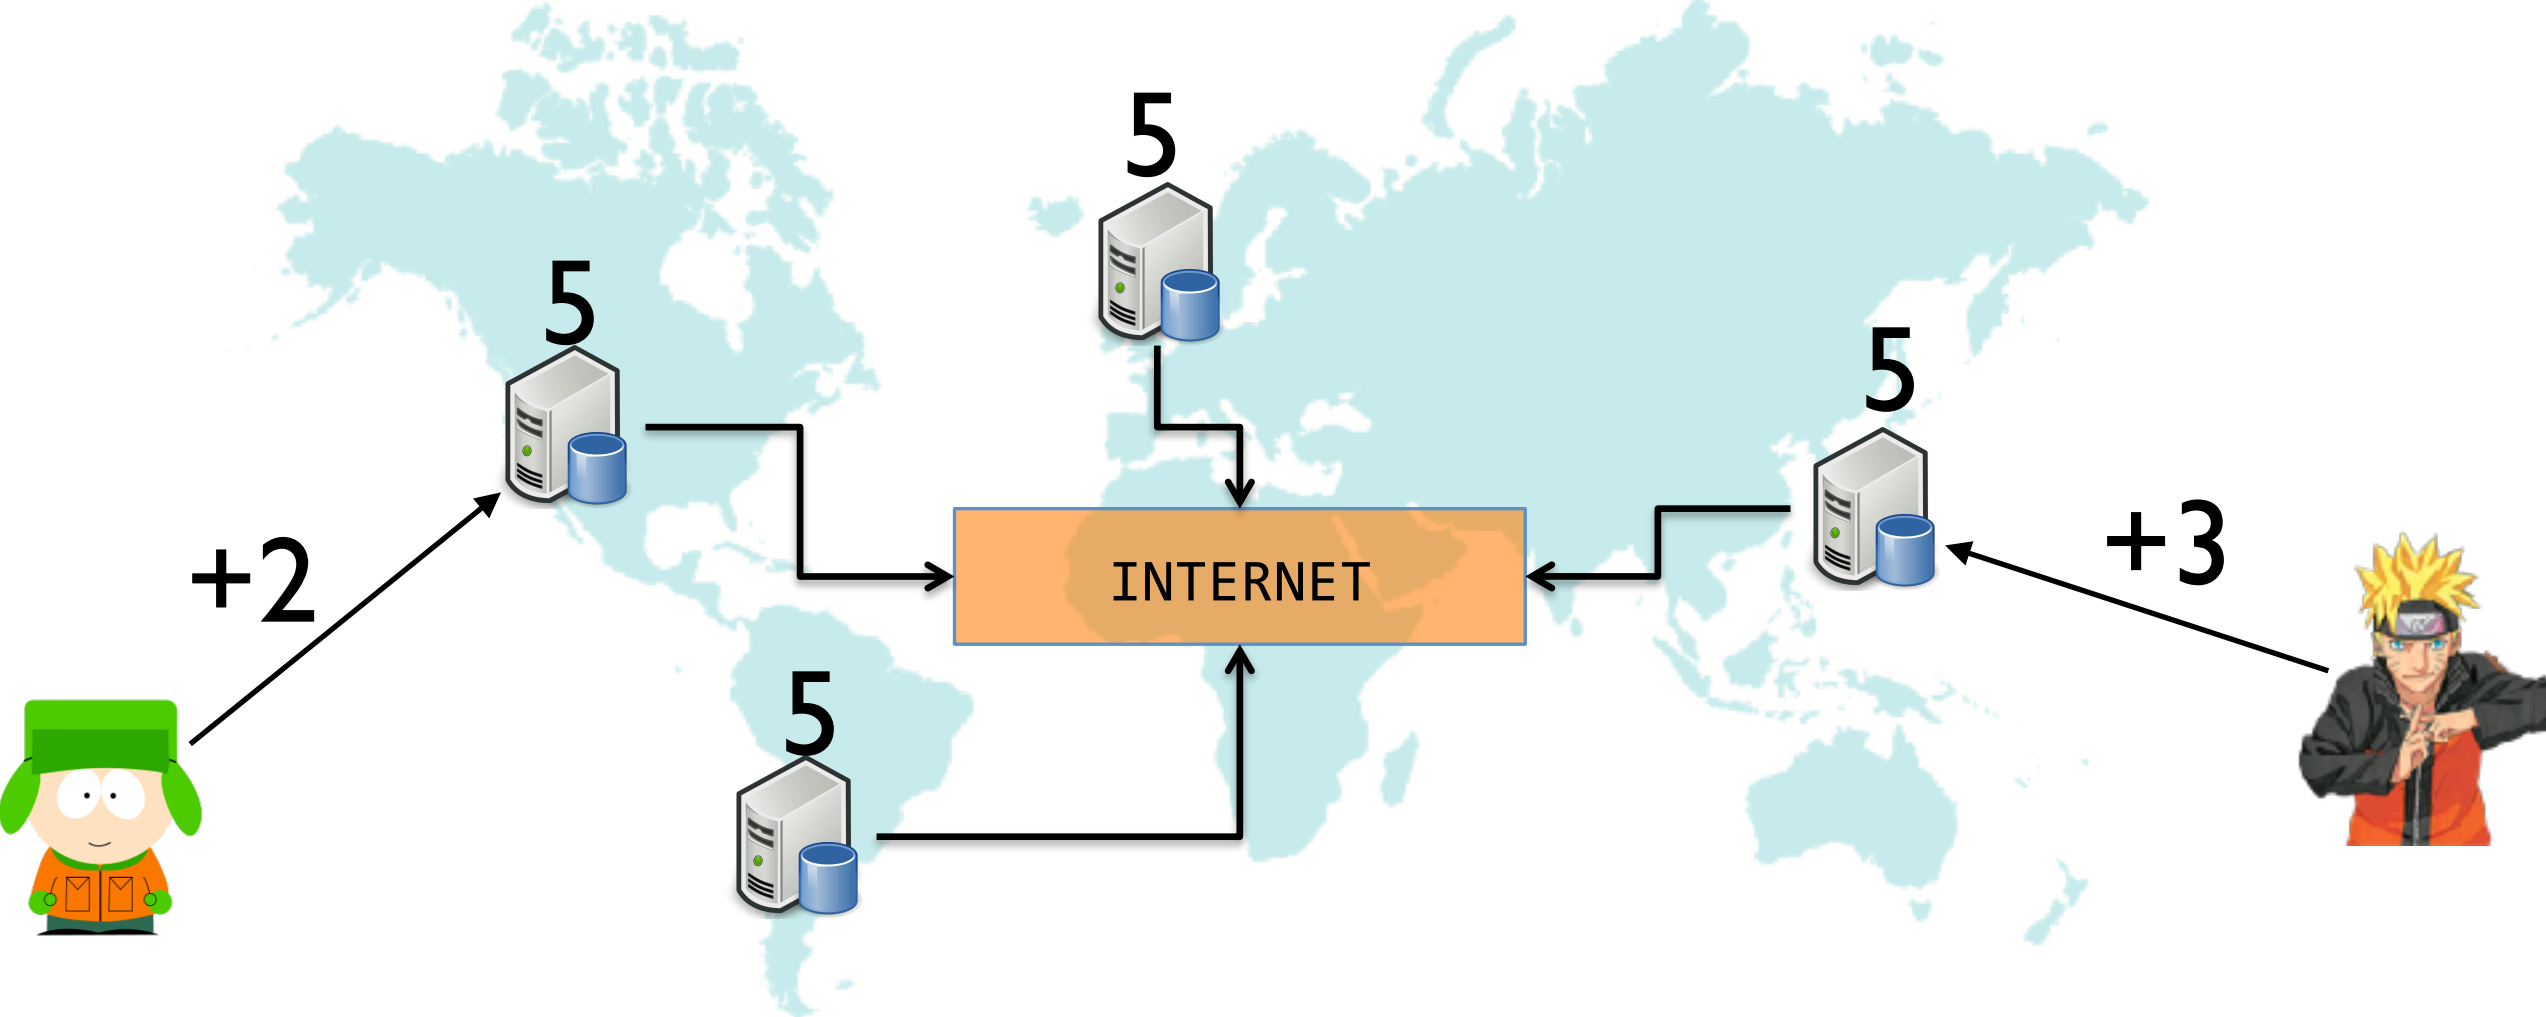
```
module Counter: sig
  type t
  val add: int → t → t
  val mult: int → t → t
  val read: t → int
end = struct
  type t = int
  let add x v = v + (abs x)
  let mult x v = v * (abs x)
  let read v = v
end
```

- Written in idiomatic functional style.
- No special reasoning principles required.
- Easily composable, e.g., `Counter.t List.t`

Replicated Counter



Replicated Counter



Adding replication to counter is non trivial!

Adding replication to counter is non trivial!

- Atomically applying updates to replicas requires expensive distributed transactions.
- Lazily applying updates leads to non-convergence since `add` and `mult` do not commute.

A common practice is to adopt the second approach, but re-engineer the datatype in terms of *commutative effects* to avoid non-convergence

Replicated Monotonic Counter Datatype in OCaml

```
module Counter: sig
  type t
  type eff
  val add: int -> t -> eff
  val mult: int -> t -> eff
  val apply: eff -> t -> t
  val read: t -> int
end = struct
  type t = int
  type eff = Add of int
  let add x v = Add (abs x)
  let mult x v = Add (v * (abs x - 1))
  let apply (Add x) v = x + v
  let read v = v
end
```

A new type for Counter effects

Counter operations return Counter effects

An apply function to apply Counter effects at local & remote replicas

Replicated Monotonic Counter Datatype in OCaml

```
module Counter: sig
  type t
  type eff
  val add: int -> t -> eff
  val mult: int -> t -> eff
  val apply: eff -> t -> t
  val read: t -> int
end = struct
```

```
  type t = int
```

```
  type eff = Add of int
```

```
  let add x v = Add (abs x)
```

```
  let mult x v = Add (v * (abs x - 1))
```

```
  let apply (Add x) v = x + v
```

```
  let read v = v
```

```
end
```

Add is the only Counter effect

The effect of mult operation is also expressed via Add

Since Add commutes with Add, counter states on all replicas eventually converge.

However ...

Replicated Counter

```
module Counter: sig
  type t
  type eff
  val add: int -> t -> eff
  val mult: int -> t -> eff
  val apply: eff -> t -> t
  val read: t -> int
end = struct
  type t = int
  type eff = Add of int
  let add x v = Add (abs x)
  let mult x v = Add (v * (abs x - 1))
  let apply (Add x) v = x + v
  let read v = v
end
```



Simple Counter

```
module Counter: sig
  type t
  val add: int -> t -> t
  val mult: int -> t -> t
  val read: t -> int
end = struct
  type t = int
  let add x v = v + (abs x)
  let mult x v = v * (abs x)
  let read v = v
end
```

Replicated Counter

```
module Counter: sig
  type t
  type eff
  val add: int -> t -> eff
  val mult: int -> t -> eff
  val apply: eff -> t -> t
  val read: t -> int
end = struct
  type t = int
  type eff = Add of int
  let add x v = Add (abs x)
  let mult x v = Add (v * (abs x - 1))
  let apply (Add x) v = x + v
  let read v = v
end
```




Simple Counter

```
module Counter: sig
  type t
  val add: int -> t -> t
  val mult: int -> t -> t
  val read: t -> int
end = struct
  type t = int
  let add x v = v + (abs x)
  let mult x v = v * (abs x)
  let read v = v
end
```

- Requires significant **re-engineering** of libraries, taking into account operational characteristics of replication (i.e., effects).
- Programming model is **restrictive**: Common polymorphic data structures (e.g., lists) are often non-commutative.
- No polymorphic structures = **no compositionality**. Each new data structure has to be engineered from scratch.

we need...

A declarative programming model for replicated functional data types.




Functional data type

+

```
module Mk_replicated :  
  functor (T:Serializable) ->  
    sig  
      ...  
    end
```


Declarative semantics of replication

=




Replicated functional data type

Functional Data Types : Persistence & Immutability




Functional Data Types



Multiple concurrent versions!

Functional Data Types \longrightarrow Functional Replicated Data Types




Multiple concurrent versions!



merge

Single replicated version

Functional Data Types \longrightarrow Functional Replicated Data Types



Functional data type


+

```
module Mk_replicated :  
  functor (T:Serializable)  
  sig  
    ...  
  end
```

Declarative semantics
of replication

\rightarrow

=



Replicated functional data type

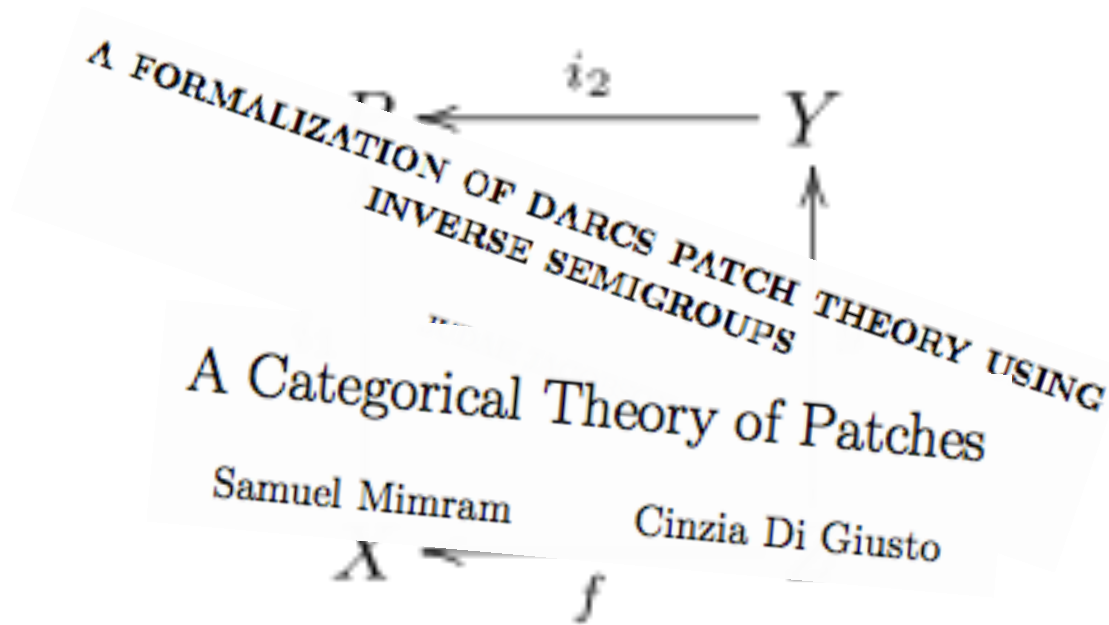
=
Semantics of merging
concurrent versions!

Mergeable Type!

Abstractly...

Merging multiple versions of data

Extensively studied for for unstructured text data



theory

&

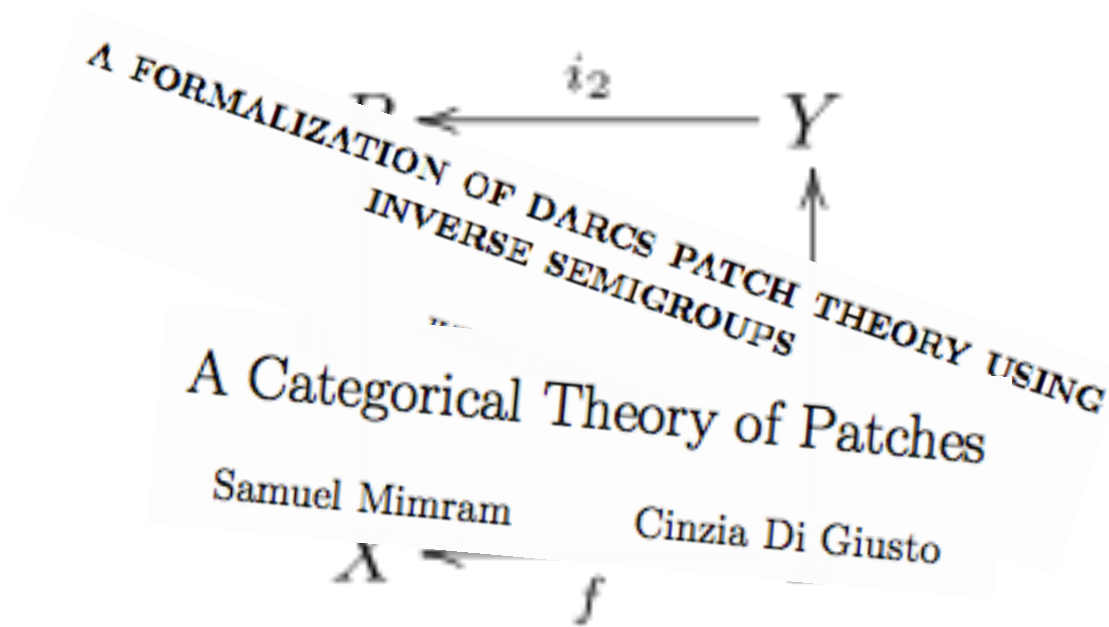
practice

of version control systems

Abstractly...

Merging multiple versions of data

Extensively studied for for unstructured text data



Mergeable Types

Extending
theory & practice
of version control systems
to functional data structures

Abstractly...

Mergeable Types

Extending
theory & practice
of version control systems
to functional data structures

Concretely...

VML:Versioned ML

An OCaml DSL that lets programmers create
version-controlled replicated data types by promoting
ordinary data types to mergeable types.

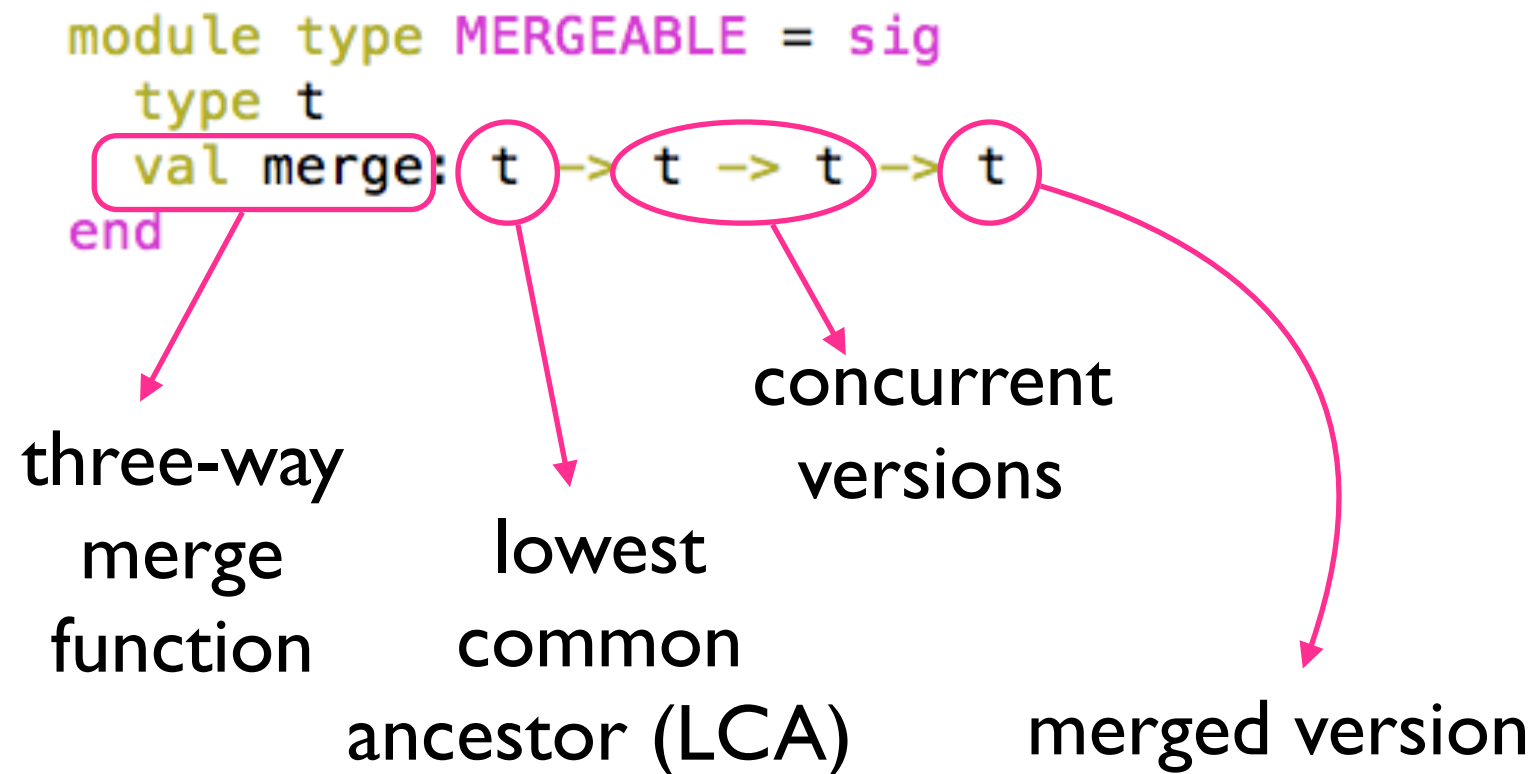
contributions

VML:Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.
2. A monad to compose concurrent/distributed computations around Versioned types.

VML:Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.

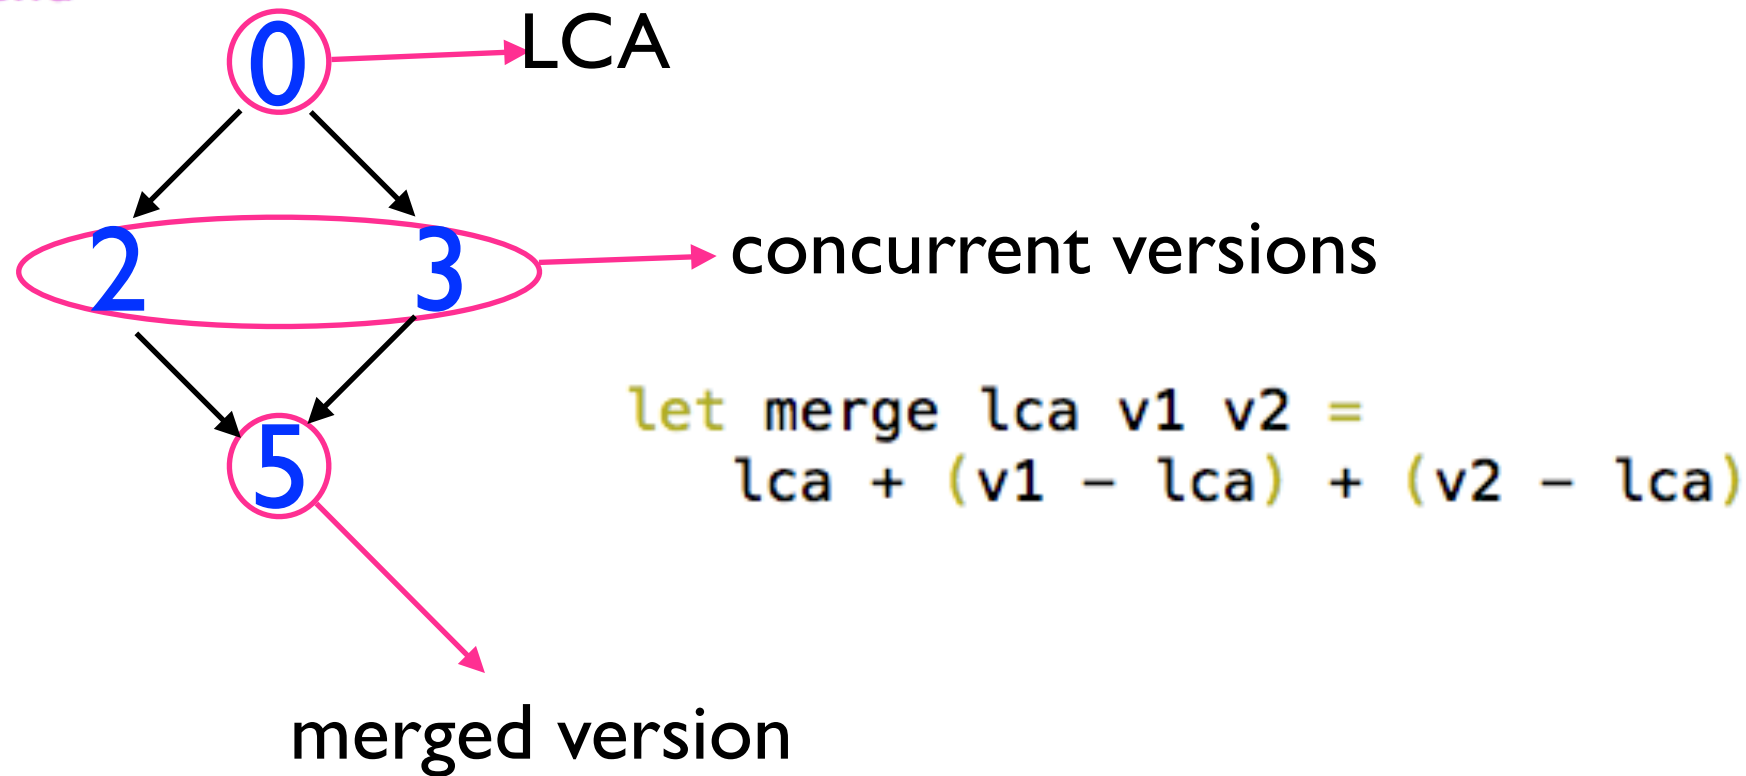


2. A monad to compose concurrent/distributed computations around Versioned types.

VML:Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.

```
module type MERGEABLE = sig
  type t
  val merge: t -> t -> t -> t
end
```



2. A monad to compose concurrent/distributed computations around Versioned types.

VML:Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.

```
module type MERGEABLE = sig
  type t
  val merge: t -> t -> t -> t
end
```



```
module type VERSIONED = sig
  type t
  type version
  val of_t: t -> version
  val to_t: version -> t
  val merge: version -> version
               -> version -> version
end
```




three-way merge
for versioned representations.


2. A monad to compose concurrent/distributed computations around Versioned types.

VML: Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.
2. A monad to compose concurrent/distributed computations around Versioned types.





VML:Versioned ML




Branching Structure

VML: Versioned ML



Mergeable Types guarantee that branches are always mergeable!

VML:Versioned ML



Mergeable Types guarantee that branches are always mergeable!*

*Conditions apply

Branches are mergeable only if a three-way merge function can be used to merge them.


```
module type MERGEABLE = sig
  type t
  val merge: t -> t -> t -> t
end
```

LCA

What if there are two LCAs?

Possible!

VML:Versioned ML



6 & 5 have two LCAs: 3 & 2!

Mergeable Types guarantee that branches are always mergeable!*

*Conditions apply

Branches are mergeable only if a three-way merge function can be used to merge them.

```
module type MERGEABLE = sig
  type t
  val merge: t -> t -> t -> t
end
```

LCA

What if there are two LCAs?

Possible!

e.g., criss-cross merges

VML:Versioned ML

1. A meta-programming framework to convert Mergeable types to Versioned types needed for replication.
2. A monad to compose concurrent/distributed computations around Versioned types.
3. A runtime that ensures progress and convergence in the presence of network partitions. **formalized**

THEOREM 4.10 (Progress). In a legal branching history H produced by the operational semantics, if two branches, b_i and b_j are not mergeable (as per Def. 4.8), then there exists a sequence of fork and merge operations (between mergeable versions) that can be performed on H to yield a new history H' where b_i and b_j are mergeable.

4. A standard library of mergeable types, including polymorphic containers, such as lists, ropes and trees. **compositionality**

VML: Versioned ML


4. A standard library of mergeable types, including polymorphic containers, such as lists, **ropes** and trees.

Usually list of chars = strings

Data structure to store a **list of values**.



Efficient concat, split, lookup, and insert at an index.

Essentially a binary tree that stores substring indexes in internal nodes, and actual strings at leaf nodes.



VML:Versioned ML

4. A standard library of mergeable types, including polymorphic containers, such as lists, **ropes** and trees.



Mergeable rope = rope + merge function

Merge boils down to merging strings

Wagner–Fischer algorithm


Wagner-Fischer + Operational Transformation

From Wikipedia, the free encyclopedia

In computer science, the **Wagner–Fischer algorithm** is a dynamic programming algorithm that computes the edit distance between two strings of characters.

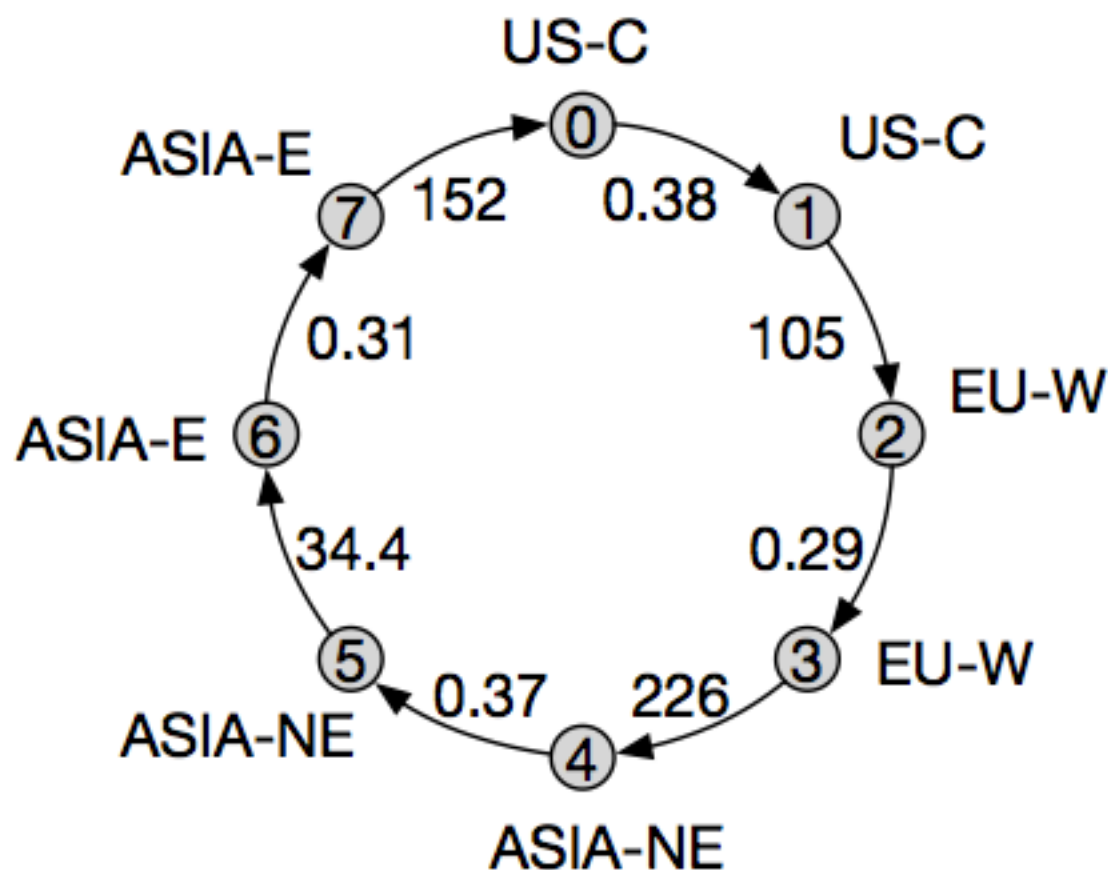
Implementation

VML is realized on top of Irmin, a persistent multi-versioned store.

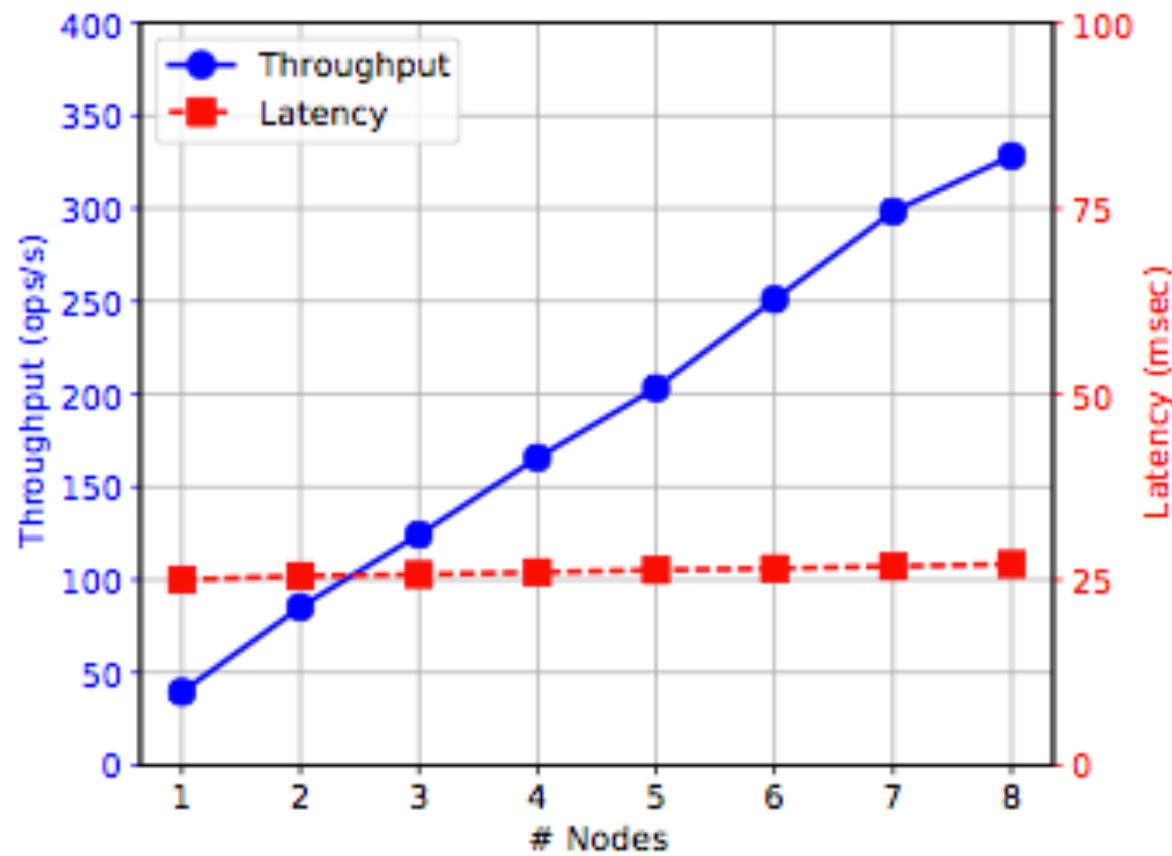
- 
- i.e., a git-like store
 - offers content-addressable heap abstraction.
address = f(content)

Evaluation

- Research Questions:
 1. Is VML practical enough to build geo-distributed eventually consistent applications?
 2. How close are we to horizontal scalability?
- Application: collaborative text editing application implemented with mergeable ropes.
- Benchmark workload: a 1576-word document. 85% insertions,
 - Setup: 8-node geo-distributed Google Compute cluster arranged in ring.
 - Each node performs edits locally, and periodically (1s) synchronizes with its successor.
 - We measured latency-throughput while increasing nodes from 1 to 8.



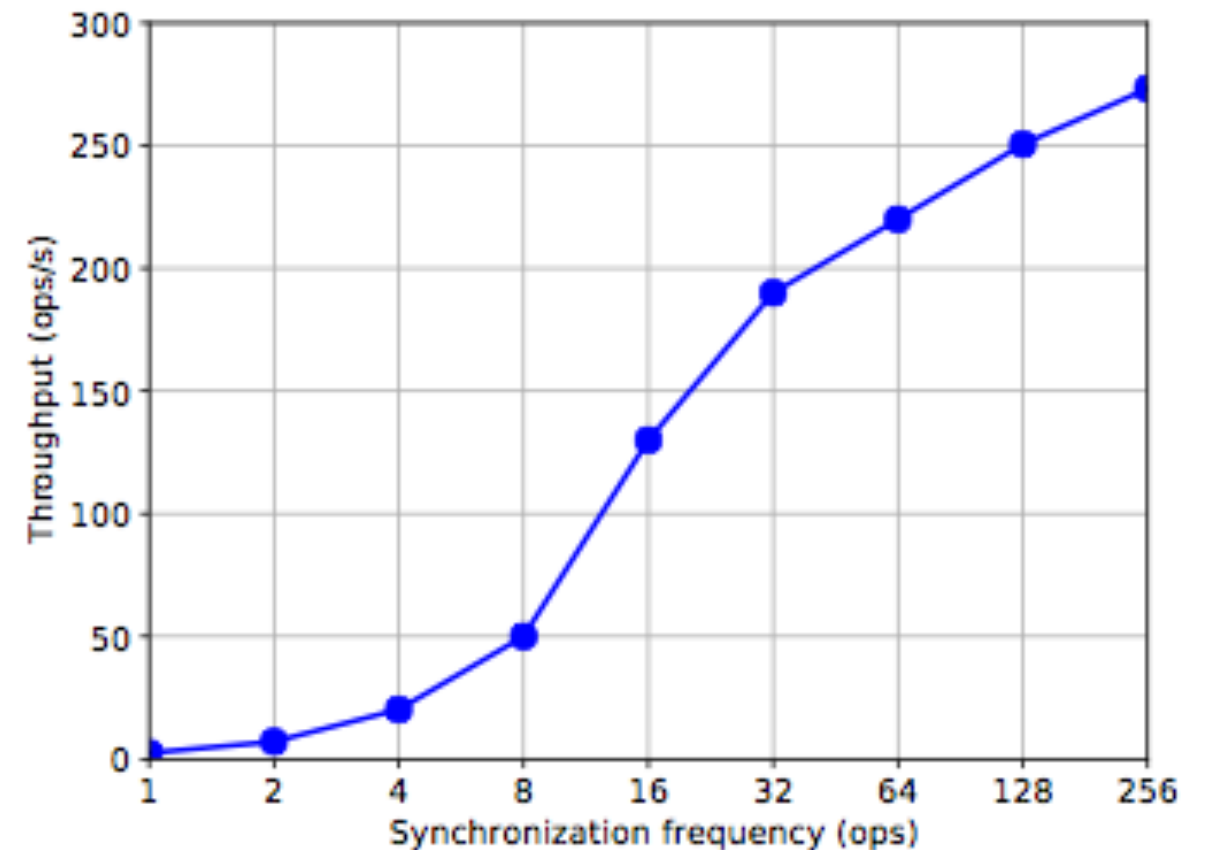
Results



(b) Scalability: Overall throughput of the cluster and latency of each operation.

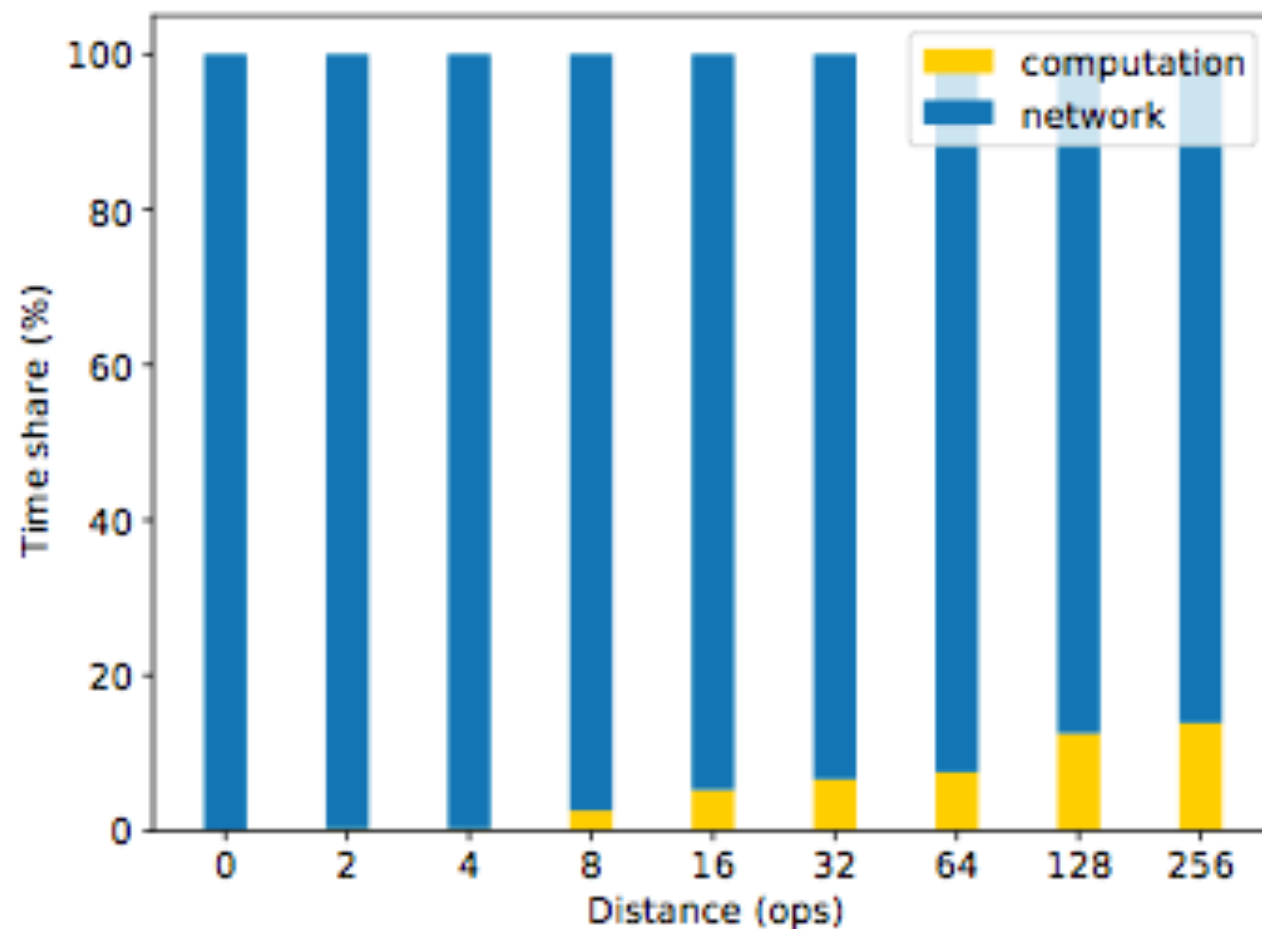
... which is not possible if we have to synchronize often.

Mergeability =>
Asynchrony =>
Near-horizontal scalability ...



(c) Synchronization: Overall throughput of the cluster while actively synchronizing with the successor node.

Results



(d) Merge performance: Cost of merging concurrent operations across nodes.

The overhead of merging is quite less when compared to the network latency.

Scope for improvement: Irmin, being a persistent store, flushes every write to the disk. Removing disk latency off the critical path should improve VML performance.

Thanks!

Short paper: <http://www.mlworkshop.org/icfp-mlworkshop17-final4.pdf>

Full paper: <http://gowthamk.github.io/docs/vml.pdf>

Code: <https://github.com/icfp2017/vml>

Replicated Counter

```
module Counter: sig
  type t
  type eff
  val add: int -> t -> eff
  val mult: int -> t -> eff
  val apply: eff -> t -> t
  val read: t -> int
end = struct
  type t = int
  type eff = Add of int
  let add x v = Add (abs x)
  let mult x v = Add (v * (abs x - 1))
  let apply (Add x) v = x + v
  let read v = v
end
```

Encodes each operation in terms
of its *effect* on the global state.

INTERNET

Is more

Lack of polymorphic types no compositionality. Each new data structure has to be engineered from scratch.

- Lack of polymorphic types no compositionality. Each new data structure has to be engineered from scratch.